

```

In [1]: #####
# Purpose: Identify a claim can be fast-tracked using Machine Learning.
# Created by: Suriya Mohan
# Created on: 12-Nov-2016
#####

# Import python - scikit-learn machine learning packages.
import numpy as np;
import pandas as pd;
from pandas import Series, DataFrame;
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import Imputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import learning_curve
from sklearn.model_selection import validation_curve
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.pipeline import Pipeline
%matplotlib inline

```

```

In [2]: # Load the data in CSV file into panda dataframe.
df_claims = pd.read_csv('C:\Users\gbu4moh\Desktop\ML - FAST TRACK\owning_adjuster

```

```

In [3]: def claim_labelEncode(df, column_name):
    ...
    Purpose: Function used to convert the nominal features (fields with string va
    Input: Dataframe and column name to encode
    Output: Modify the passed dataframe by adding encoded field to end and drop t
    ...
    df[column_name].fillna(value=df[column_name].value_counts().idxmax(), inplace

    le = LabelEncoder()
    encoder = le.fit_transform(df[column_name].values)
    df.insert(0, 'code' + column_name, encoder.astype(int))
    df.drop(column_name, axis=1, inplace=True)

```

```
In [4]: def claim_oneHotEncode(df,column_array):
        ...
        Purpose: Function used to hot encode fields - example field having 3 distinct
        Input: Dataframe and column array
        Output: Returns the data frame with encoded values.
        ...
        ohe = OneHotEncoder(categorical_features=column_array,sparse=False)
        df_ohe = ohe.fit_transform(df)
        return df_ohe
```

```
In [5]: def claim_standardScaler(df):
        ...
        Purpose: Normalize the dataframe so the features are on the same scale.
        Input: Dataframe
        Output: Returns the data frame with standarized values.
        ...
        sc = StandardScaler()
        df_std = sc.fit_transform(df)
        return df_std
```

```
In [6]: def claim_Imputer(df):
        ...
        Purpose: Normalize the dataframe so the features are on the same scale.
        Input: Dataframe
        Output: Returns the data frame with standarized values.
        ...
        imr = Imputer(missing_values=np.NaN, strategy='most_frequent',axis=0)
        imr = imr.fit_transform(df)
        return imr
```

```
In [7]: # Display all the fields in the input dataframe.
        df_claims.columns
```

```
Out[7]: Index([u'CLAIM_NUM', u'CASUALTY_UNIT_IND', u'FATALITY_IND',
               u'HIGH_VALUE_CLAIM_IND', u'ACCIDENT_FAULT_IND', u'LEGAL_ENTITY_NM',
               u'WEATHER_CD_DESC', u'CLAIM_JURISDICTION_ST_CD_DESC',
               u'LOSS_TYPE_CD_DESC', u'CLAIM_COMPLEXITY_CD_DESC',
               u'CLAIM_FILE_TYPE_DESC', u'CLAIM_TIER_CD_DESC', u'FAULT_RATING_CD_DESC',
               u'CATASTROPHE_IND', u'LOSS_SUB_TYPE_DESC', u'LIABILITY_CD_DESC',
               u'LIABILITY_PERCENTAGE', u'INCIDENT_ONLY_IND', u'GLASS_ONLY_IND',
               u'NON_CHARGEABLE_CLAIM_IND', u'TOTAL_LOSS_IND', u'PRODUCT_TYPE_DESC',
               u'ATTORNEY_ON_CLAIM_IND', u'VEHICLE_TOW_STORAGE_IND',
               u'COVERAGE_TYPE_CD_DESC', u'TOW_TYPE_CD_DESC', u'INCIDENT_TYPE_CD_DESC',
               u'COVERAGE_SUB_TYPE_CD_DESC', u'CLAIM_TYPE_DESC', u'COVERAGE_NM',
               u'PROPERTY_DESC', u'DAMAGE_DESC', u'VEH_COUNT',
               u'CLAIM_ADJUSTER_GROUP_TYPE'],
              dtype='object')
```

```
In [1]: # Check value and count of each value in all the fields in the dataframe.
# for col in df_claims.columns:
#     if col in ['CLAIM_NUM']:
#         continue

#     print col
#     value_cnt = df_claims[col].value_counts()
#     print value_cnt
#     print '*****'
```

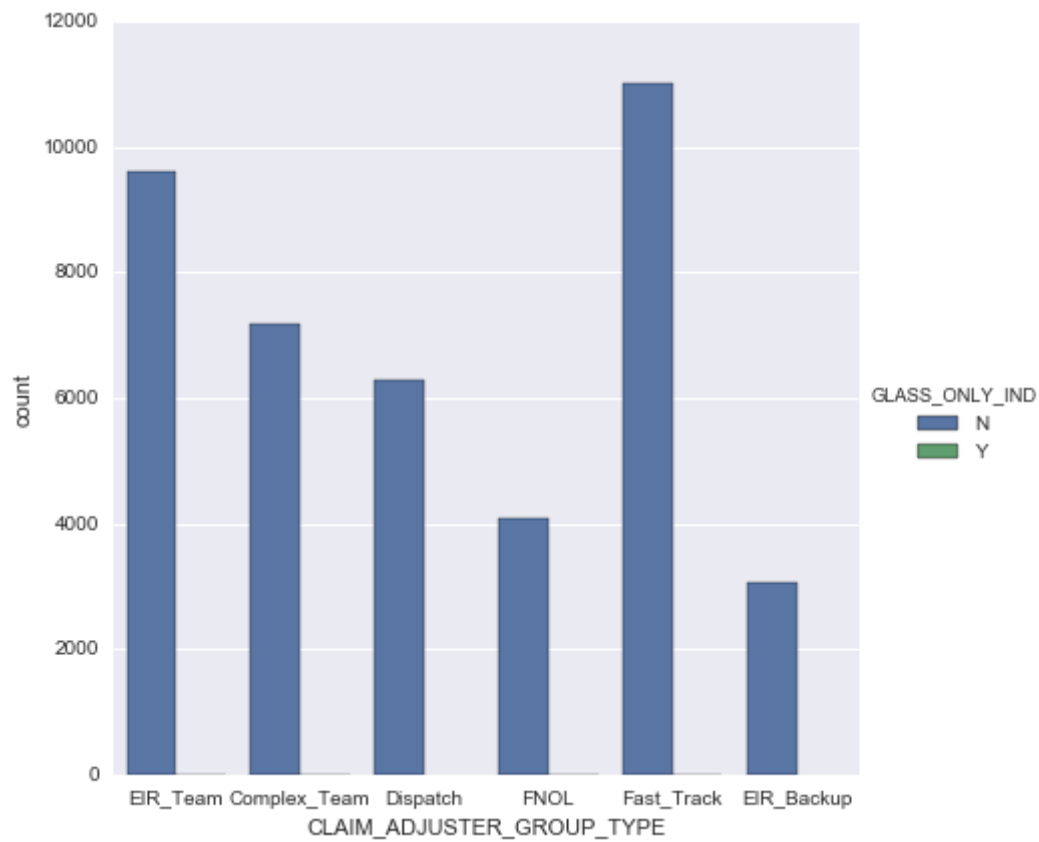
```
In [10]: # Move the target field to array.['CLAIM_NUM', 'CLAIM_ADJUSTER_GROUP_TYPE']
df_target = df_claims[['CLAIM_NUM', 'CLAIM_ADJUSTER_GROUP_TYPE']]

df_target.groupby('CLAIM_ADJUSTER_GROUP_TYPE').count()
```

Out[10]:


	CLAIM_NUM
CLAIM_ADJUSTER_GROUP_TYPE	
Complex_Team	7181
Dispatch	6297
EIR_Backup	3047
EIR_Team	9616
FNOL	4089
Fast_Track	11022

```
In [11]: # Plot the glass only indicator and claim adjuster group.  
df_report = df_claims[df_claims['CLAIM_ADJUSTER_GROUP_TYPE'] != 'FNOL']  
  
sns.factorplot('CLAIM_ADJUSTER_GROUP_TYPE', data=df_claims, hue='GLASS_ONLY_IND', k  
Out[11]: <seaborn.axisgrid.FacetGrid at 0x3f43c50>
```



```
In [12]: # Plot the fault rating code and claim adjuster group.  
df_report = df_claims[df_claims['CLAIM_ADJUSTER_GROUP_TYPE'] != 'FNOL']  
  
sns.factorplot('CLAIM_ADJUSTER_GROUP_TYPE', data=df_claims, hue='FAULT_RATING_CD_DE
```


```
Out[12]: <seaborn.axisgrid.FacetGrid at 0xd555e80>
```



```
In [13]: # Plot the coverage code and claim adjuster group.
df_report = df_claims[df_claims['CLAIM_ADJUSTER_GROUP_TYPE'] != 'FNOL']

sns.factorplot('CLAIM_ADJUSTER_GROUP_TYPE', data=df_claims, hue='COVERAGE_TYPE_CD_D
```

```
Out[13]: <seaborn.axisgrid.FacetGrid at 0xd555f60>
```



```
In [14]: # drop fields which are not useful -
df_claims.drop(['DAMAGE_DESC', 'PROPERTY_DESC', 'CLAIM_TYPE_DESC', 'INCIDENT_TYPE_',
'NON_CHARGEABLE_CLAIM_IND', 'INCIDENT_ONLY_IND', 'LOSS_SUB_TYPE_DESC', 'CLAIM_FILE_
```

```
In [15]: # Encode the nominal features - String to integers.
for col in df_claims.columns:
    if col in ['VEH_COUNT', 'LIABILITY_PERCENTAGE']:
        continue

    claim_labelEncode(df_claims, col)
```

```
In [16]: # Impute the missing values in the dataframe.
claim_imp = claim_Imputer(df_claims)
```


```
In [17]: # Convert to dataframe.
df_claim_imp = pd.DataFrame(data=claim_imp, columns=df_claims.columns)
```

```
In [18]: # Standarize the features so the features are on same - scale.  
# array_std = claim_standardScaler(df_claim_imp[['LIABILITY_PERCENTAGE']])  
  
# df_std = pd.DataFrame(data=array_std,columns=['LIABILITY_PERCENTAGE'])  
  
# df_claim_imp.drop(['LIABILITY_PERCENTAGE'], inplace=True,axis=1)  
  
# df_claim_std = pd.concat([df_claim_imp,df_std],axis=1)  
  
df_claim_std = df_claim_imp  
df_claim_std.head()
```

Out[18]:

	codeCOVERAGE_NM	codeCOVERAGE_SUB_TYPE_CD_DESC	codeTOW_TYPE_CD_DE!
0	1.0	0.0	4.0
1	5.0	5.0	4.0
2	6.0	6.0	4.0
3	1.0	0.0	4.0
4	5.0	5.0	4.0

5 rows × 21 columns



In [19]: *# Create target dataframe. Set the fast track values to 1 and other groups to 0.*
`df_target.loc[df_target['CLAIM_ADJUSTER_GROUP_TYPE'] != 'Fast_Track', 'CLAIM_ADJU
df_target.loc[df_target['CLAIM_ADJUSTER_GROUP_TYPE'] == 'Fast_Track', 'CLAIM_ADJU
df_target.groupby('CLAIM_ADJUSTER_GROUP_TYPE').count()`

C:\Users\gbu4moh\AppData\Local\Enthought\Canopy\User\lib\site-packages\pandas\core\indexing.py:128: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

`self._setitem_with_indexer(indexer, value)`

C:\Users\gbu4moh\AppData\Local\Enthought\Canopy\User\lib\site-packages\ipykernel__main__.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

`from ipykernel import kernelapp as app`

C:\Users\gbu4moh\AppData\Local\Enthought\Canopy\User\lib\site-packages\ipykernel__main__.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

`app.launch_new_instance()`

Out[19]:

	CLAIM_NUM
CLAIM_ADJUSTER_GROUP_TYPE	
0	30230
1	11022


```
In [20]: # Use the Random forest classifier to identify the variable importance.
forest = RandomForestClassifier(n_estimators=1000, random_state=0,n_jobs=-1)

X_train, X_test, y_train, y_test = train_test_split(df_claim_std, list(df_target[


labels = df_claim_std.columns
forest.fit(X_train,y_train)
importances = forest.feature_importances_
indices = np.argsort(importances) [::-1]

for f in range(X_train.shape[1]):
    print(labels[f], importances[indices[f]])

feat_imp = pd.Series(importances, labels).sort_values(ascending=False)
feat_imp.plot(kind='bar', title='Feature Importances',figsize=(10,10))
plt.ylabel('Feature Importance Score')
```

```
('codeCOVERAGE_NM', 0.33120909816568078)
('codeCOVERAGE_SUB_TYPE_CD_DESC', 0.32737019119593369)
('codeTOW_TYPE_CD_DESC', 0.095948524702808485)
('codeCOVERAGE_TYPE_CD_DESC', 0.052279342014457809)
('codeVEHICLE_TOW_STORAGE_IND', 0.040380598935093506)
('codeATTORNEY_ON_CLAIM_IND', 0.023023529470336007)
('codePRODUCT_TYPE_DESC', 0.021875706496223352)
('codeGLASS_ONLY_IND', 0.019732626939262192)
('codeLIABILITY_CD_DESC', 0.019581345394470415)
('codeCATASTROPHE_IND', 0.019073826759692811)
('codeFAULT_RATING_CD_DESC', 0.016952099438646343)
('codeCLAIM_TIER_CD_DESC', 0.01566834599296309)
('codeCLAIM_COMPLEXITY_CD_DESC', 0.0079588240237430802)
('codeLOSS_TYPE_CD_DESC', 0.0060141628108448435)
('codeCLAIM_JURISDICTION_ST_CD_DESC', 0.0019210044761429716)
('codeWEATHER_CD_DESC', 0.00076214346082284555)
('codeLEGAL_ENTITY_NM', 0.0002486297228772104)
('codeACCIDENT_FAULT_IND', 0.0)
('codeHIGH_VALUE_CLAIM_IND', 0.0)
('codeFATALITY_IND', 0.0)
('codeCASUALTY_UNIT_IND', 0.0)
```

```
Out[20]: <matplotlib.text.Text at 0xd8d3b00>
```



```
In [21]: # first 23 fields needs to be hot-encoded.
ohe = OneHotEncoder(categorical_features=np.arange(0,21))
df_claim_hot = ohe.fit_transform(df_claim_std)
```

```
In [22]: # Create data frame of hot-encoded array.
df_claim_mod = pd.DataFrame(df_claim_hot.toarray())
```

```
In [23]: # Split the dataset to train and test dataset. 70% of data is trained and 30% of
X_train, X_test, y_train, y_test = train_test_split(df_claim_mod, list(df_target[
```

```
In [24]: # Train the Logistic Regression model.
param_range = [0.01,0.02,0.03,0.04,0.05,0.06,0.07,0.08,0.09]
```

```
for param in param_range:
    print('C : ',param)
    lr = LogisticRegression(penalty='l2',C=param)
    lr.fit(X_train,y_train)
    print('training accuracy:', lr.score(X_train,y_train))
    print('test accuracy: ', lr.score(X_test, y_test))
```

```
('C : ', 0.01)
('training accuracy:', 0.77392990718936139)
('test accuracy: ', 0.76656431803490632)
('C : ', 0.02)
('training accuracy:', 0.7774622523895276)
('test accuracy: ', 0.76979638009049778)
('C : ', 0.03)
('training accuracy:', 0.78078681257791938)
('test accuracy: ', 0.77302844214608923)
('C : ', 0.04)
('training accuracy:', 0.78109849009558108)
('test accuracy: ', 0.7738364576599871)
('C : ', 0.05)
('training accuracy:', 0.78175647596620035)
('test accuracy: ', 0.77674531351001941)
('C : ', 0.06)
('training accuracy:', 0.78186036847208751)
('test accuracy: ', 0.77674531351001941)
('C : ', 0.07)
('training accuracy:', 0.78196426097797478)
('test accuracy: ', 0.77682611506140919)
('C : ', 0.08)
('training accuracy:', 0.78203352264856629)
('test accuracy: ', 0.77682611506140919)
('C : ', 0.09)
('training accuracy:', 0.78213741515445356)
('test accuracy: ', 0.77690691661279898)
```

```
In [25]: # Train the SVM model.
param_range = [0.0001,0.001,0.01,0.1,1.0,10.0,100.0,1000.0]

for param in param_range:
    svm = SVC(kernel='sigmoid',C=param,random_state= 0)
    svm.fit(X_train,y_train)

    print('training accuracy:', svm.score(X_train,y_train))
    print('test accuracy: ', svm.score(X_test, y_test))

('training accuracy:', 0.73566283418756062)
('test accuracy: ', 0.7261635423400129)
('training accuracy:', 0.73566283418756062)
('test accuracy: ', 0.7261635423400129)
('training accuracy:', 0.73566283418756062)
('test accuracy: ', 0.7261635423400129)
('training accuracy:', 0.74993073832940849)
('test accuracy: ', 0.73965740142210734)
('training accuracy:', 0.76856212771852062)
('test accuracy: ', 0.75816095669036843)
('training accuracy:', 0.7746225238952763)
('test accuracy: ', 0.76454427925016155)
('training accuracy:', 0.73074525557556447)
('test accuracy: ', 0.72745636716224948)
('training accuracy:', 0.73922981022302259)
('test accuracy: ', 0.73553652230122823)
```

```

In [26]: # Plot the training Curve
pipe_lr = Pipeline([('clf', LogisticRegression(penalty='l2', random_state=0, C=0.01)

train_sizes, train_scores, test_scores = learning_curve(estimator=pipe_lr,
                                                         X=X_train,
                                                         y=y_train,
                                                         train_sizes=np.linspace(0.1, 1.0, 10),
                                                         cv=10,
                                                         n_jobs=1
                                                         )

# for x,y,z in zip(train_sizes, train_scores, test_scores):
#     print('train size ', x)
#     print('train scores ', y)
#     print('test scores ', z)

train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

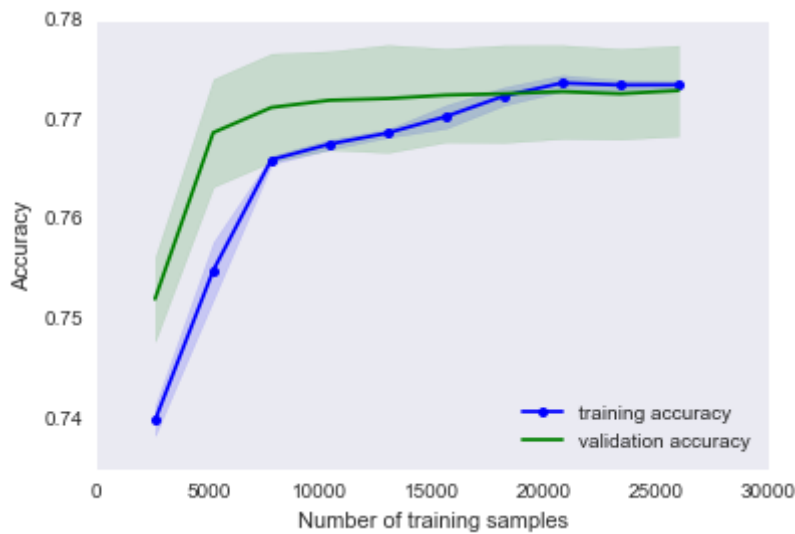
print('train_mean ', train_mean)
print('train_std ', train_std)
print('test_mean ', test_mean)
print('test_std ', test_std)

plt.plot(train_sizes, train_mean, color='blue', marker='o', markersize=5, label='train
plt.fill_between(train_sizes, train_mean + train_std, train_mean - train_std, alp
plt.plot(train_sizes, test_mean, color='green', marker='x', markersize=5, label='valid
plt.fill_between(train_sizes, test_mean + test_std, test_mean - test_std, alpha=0

plt.grid()
plt.xlabel('Number of training samples')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
# plt.ylim([0.8, 1.0])
plt.show()

('train_mean ', array([ 0.73991532,  0.75495478,  0.76608517,  0.76765442,  0.7
687678 ,
                    0.77042073,  0.77249038,  0.77379864,  0.7735933 ,  0.77360988]))
('train_std ', array([ 0.00150115,  0.00289429,  0.00042232,  0.00051272,  0.00
041475,
                    0.00117043,  0.00090807,  0.0008127 ,  0.00053942,  0.00042531]))
('test_mean ', array([ 0.75211247,  0.76880529,  0.77133337,  0.77206079,  0.77
223383,
                    0.77258009,  0.77271865,  0.7729264 ,  0.77271855,  0.77303008]))
('test_std ', array([ 0.00422554,  0.00541723,  0.0054412 ,  0.00495734,  0.005
41199,
                    0.00470505,  0.00489368,  0.00470259,  0.00455039,  0.00455052]))

```



```
In [27]: # confusion matrix

y_predict = lr.predict(X_test)

confusion_matrix(y_test, y_predict)

#print ("Accuracy : %.4g" % metrics.accuracy_score(y_test, y_predict))
#print ("AUC Score (Train): %f" % metrics.roc_auc_score(y_test, y_predict))

#cv_score = cross_validation.cross_val_score(alg, dtrain[predictors], dtrain['Dis
# print "CV Score : Mean - %.7g | Std - %.7g | Min - %.7g | Max - %.7g" % (np.me
```

```
Out[27]: array([[8103,  884],
               [1877, 1512]])
```

```
In [28]: # Gradient boosting classifier.
param_test = {'n_estimators':range(20,81,10)}
gsearch1 = GridSearchCV(estimator = GradientBoostingClassifier(learning_rate=0.1,
param_grid = param_test, scoring='roc_auc',n_jobs=4,iid=False, cv=5)
gsearch1.fit(X_train,y_train)
```

```
Out[28]: GridSearchCV(cv=5, error_score='raise',
                    estimator=GradientBoostingClassifier(criterion='friedman_mse', init=None,
                    learning_rate=0.1, loss='deviance', max_depth=8,
                    max_features='sqrt', max_leaf_nodes=None,
                    min_impurity_split=1e-07, min_samples_leaf=50,
                    min_samples_split=500, min_weight_fraction_leaf=0.0,
                    n_estimators=100, presort='auto', random_state=10,
                    subsample=0.8, verbose=0, warm_start=False),
                    fit_params={}, iid=False, n_jobs=4,
                    param_grid={'n_estimators': [20, 30, 40, 50, 60, 70, 80]},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='roc_auc', verbose=0)
```

```
In [29]: gsearch1.grid_scores_, gsearch1.best_params_, gsearch1.best_score_
```

```
C:\Users\gbu4moh\AppData\Local\Enthought\Canopy\User\lib\site-packages\sklearn  
\model_selection\_search.py:662: DeprecationWarning: The grid_scores_ attribute  
was deprecated in version 0.18 in favor of the more elaborate cv_results_ attr  
ibute. The grid_scores_ attribute will not be available from 0.20  
DeprecationWarning)
```

```
Out[29]: ([mean: 0.84697, std: 0.00728, params: {'n_estimators': 20},  
mean: 0.84918, std: 0.00776, params: {'n_estimators': 30},  
mean: 0.85005, std: 0.00813, params: {'n_estimators': 40},  
mean: 0.85093, std: 0.00823, params: {'n_estimators': 50},  
mean: 0.85120, std: 0.00795, params: {'n_estimators': 60},  
mean: 0.85174, std: 0.00808, params: {'n_estimators': 70},  
mean: 0.85182, std: 0.00810, params: {'n_estimators': 80}],  
{'n_estimators': 80},  
0.85181943328741738)
```

```
In [30]: y_predict = gsearch1.predict(X_test)  
  
confusion_matrix(y_test, y_predict)
```

```
Out[30]: array([[7591, 1396],  
[1233, 2156]])
```

```
In [31]: print('training accuracy:', gsearch1.score(X_train,y_train))  
print('test accuracy: ', gsearch1.score(X_test, y_test))  
  
( 'training accuracy:', 0.85568745454417727)  
( 'test accuracy: ', 0.85105734347665818)
```

```
In [35]: import theano
from theano import tensor

# from keras.models import Sequential
# from keras.layers.core import Dense
# from keras.optimizers import SGD

# np.random.seed(1)

# model = Sequential()
# model.add(Dense(input_dim=X_train.shape[1],
#                  output_dim=50,
#                  init='uniform',
#                  activation='tanh'))

# model.add(Dense(input_dim=50,
#                  output_dim=50,
#                  init='uniform',
#                  activation='tanh'))

# model.add(Dense(input_dim=50,
#                  output_dim=y_train_ohe.shape[1],
#                  init='uniform',
#                  activation='softmax'))

# sgd = SGD(lr=0.001, decay=1e-7, momentum=.9)
# model.compile(loss='categorical_crossentropy', optimizer=sgd)

# model.fit(X_train, y_train,
#           nb_epoch=50,
#           batch_size=300,
#           verbose=1,
#           validation_split=0.1,
#           show_accuracy=True)
```

```
In [ ]:
```